

BioSys PhD | Earthsystems PhD

Statistics 1

Exploratory Data Analysis

Lisete Sousa

lmsousa@fc.ul.pt

Room: 6.4.25

*Departamento de Estatística e Investigação Operacional
Faculdade de Ciências da Universidade de Lisboa*

2017

Summary

1. Descriptive Statistics
2. Graphical Representations
 - Bar plot
 - Circular Diagram
 - Histogram
 - Box plot
 - Quantile-Quantile Plot

1. Descriptive Statistics

R, through package `stats`, allows the calculation of sample characteristics for one, or more, variables simultaneously.

- `mean()` → arithmetic mean
- `median()` → median
- `sd()` → standard deviation
- `var()` → variance
- `quantile()` → quantiles
- `min()` → minimum
- `max()` → maximum
- ...

Example:

```
> data<-c(2,5,3,7,1,4,2,5,2,7,1,3,2,3,5,6)
> mean(data)
[1] 3.625
> median(data)
[1] 3
> var(data)
[1] 3.983333
> sd(data)
[1] 1.995829
> quantile(data)

0%   25%   50%   75%  100%
1    2    3    5    7

> min(data)
[1] 1
> max(data)
[1] 7
```

Notes:

- 1 We may check the standard deviation by doing:

```
> sqrt(var(data))
[1] 1.995829
```

- 2 It is possible to calculate other quantiles by using function `quantile()`:

```
quantile(data,c(0.2,0.7))
20%    70%
2      5
```

Let us calculate some of these characteristics for the database `juul`, available through package `ISwR`, starting with assessing the variables contained therein.

```
>names(juul)
[1] "age" "menarche" "sex" "igf1" "tanner" "testvol"
```

To obtain the average ages we have to do:

```
>attach(juul)
>mean(age)
[1] NA
```

The average ages is not available (NA) because there are missing values

For these NA values to be ignored we have to include an extra argument:

```
> mean(age,na.rm=T)
[1] 15.09535

> colMeans(juul,na.rm=T)
```

age	menarche	sex	igf1	tanner	testvol
15.095352	1.475852	1.534483	340.167976	2.639672	7.895833

Note that the values indicated by the color red should not be considered because they are related to qualitative variables.

summary()

Function `summary` allows us to view a set of sample features concerning a certain variable.

```
> summary(age)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
0.170	9.053	12.560	15.100	16.860	83.000	5.000

This function also returns the number of missing data.

This same function allows us to obtain a summary for all variables in the database:

```
> summary(juul)
```

```
> summary(juul)
```

age	menarche	sex	igf1	tanner	testvol
Min. : 0.170	Min. : 1.000	Min. :1.000	Min. : 25.0	Min. : 1.000	Min. : 1.000
1st Qu.: 9.053	1st Qu.: 1.000	1st Qu.:1.000	1st Qu.:202.3	1st Qu.: 1.000	1st Qu.: 1.000
Median :12.560	Median : 1.000	Median :2.000	Median :313.5	Median : 2.000	Median : 3.000
Mean :15.095	Mean : 1.476	Mean :1.534	Mean :340.2	Mean : 2.640	Mean : 7.896
3rd Qu.:16.855	3rd Qu.: 2.000	3rd Qu.:2.000	3rd Qu.:462.8	3rd Qu.: 5.000	3rd Qu.: 15.000
Max. :83.000	Max. : 2.000	Max. :2.000	Max. :915.0	Max. : 5.000	Max. : 30.000
NA's : 5.000	NA's :635.000	NA's :5.000	NA's :321.0	NA's :240.000	NA's :859.000

NOTE: The qualitative variables (sex, menarche and tanner) are encoded as numerical (quantitative). This has to be changed as follows:

```
> juul$sex<-factor(juul$sex, labels=c("M","F"))  
> juul$menarche<-factor(juul$menarche, labels=c("No","Yes"))  
> juul$tanner<-factor(juul$tanner, labels=c("I","II","III",  
+ "IV","V"))
```

Exercise:

Check now the new form of the summary table.

tapply()

Function `tapply()` calculates sample characteristics for the given variable for each level of a second variable, which may be qualitative (or categorical).

Example:

The average of *igf1* per gender:

```
> tapply(igf1,sex,mean,na.rm=T)
      M      F
310.8866 368.1006
```

NOTA: The argument `na.rm=T` must be inserted due to the presence of missing data.

table()

A simple way to describe qualitative data (and also grouped data) consists of building tables.

For univariate table construction the procedure is simple:

```
> table(sex)
sex
  M    F
621 3713
```

For two variables a double-entry table is implemented:

```
> table(sex,tanner)
```

	tanner				
sex	I	II	III	IV	V
M	291	55	34	41	124
F	224	48	38	40	204

Is is also possible to build a triple-entry table:

```
> table(menarche,sex,tanner)
```

Let us see now, how to use R, trough the package **graphics**, to do graphical representations.

NOTE: There are other packages with graphical functions in R, such as, **gplots**, **lattice**, **misc3d**, etc.

2. Graphical Representations

Bar Plot

Bar Plot

The bar plot (diagram of bars) is built using the function `barplot()`.

If we want to represent variable `tanner` through a bar plot (absolute frequencies), we do:

```
> barplot(table(tanner))
```


If we want to represent tanner in a bar plot, with relative frequencies, we do:


```
> table(tanner)/sum(!is.na(tanner))
```

NOTE: The relative frequency must be calculated based on the total number of non-missing observations. We run the function `sum(!is.na(tanner))`, or `length(tanner)-sum(is.na(tanner))`, to calculate the number of non-missing values.

The graphic may be improved if we include some more arguments:

```
> barplot(table(tanner)/(length(tanner)-240),xlab="tanner",  
+ ylab="Relative Freq.",ylim=c(0,0.5))
```

```
> par(mfrow=c(1,3))  
> barplot(table(tanner))  
> sum(is.na(tanner))  
> table(tanner)/(length(tanner)-240)  
> barplot(table(tanner)/(length(tanner)-240))  
> barplot(table(tanner)/(length(tanner)-240),xlab="tanner",  
+ ylab="Relative Freq.",ylim=c(0,0.5))
```



Circular Diagram

Circular Diagram (Pie Chart)

The circular diagram is built using the function `pie()`.



This diagram is valid for qualitative data and grouped data.

```
> pie(table(tanner))
```

We can change the colors and insert a title:

```
> slices<-c("yellow","orange","red","blue","green")  
> pie(table(tanner),col=slices,main="tanner")
```

```
> par(mfrow=c(1,2),mex=0.3,mar=c(1,0))  
> pie(table(tanner))  
> slices<-c("yellow","orange","red","blue","green")  
> pie(table(tanner),col=slices,main="tanner")
```



Histogram

Histogram

The histogram is built using the function `hist()`.

```
> hist(igf1)
```

By default, this function uses the rule of Sturges to build classes and represents the absolute frequencies.

You can also view the results in the form of text:

```
> hist(igf1,plot=FALSE)
```

You may be interested in:

- color the bars and represent the *densities* (relative frequency divided by the range of each class)

```
> hist(igf1,probability=TRUE,col="blue")
```

- define other classes

```
> class<-c(0,50,200,350,500,650,1000)
```

- put the absolute frequency at the top of each bar and shade them

```
> h<-hist(igf1,breaks=class,angle=45,density=40)
> h
> text(h$mids,h$density,h$counts,adj=c(0.5,-1))
```

```
> hist(igf1,plot=FALSE)
$breaks
 [1]    0  100  200  300  400  500  600  700  800  900 1000

$counts
 [1]  43 204 247 163 185  98  43  24   9   2




$intensities
 [1] 4.223968e-04 2.003929e-03 2.426326e-03 1.601179e-03 1.817289e-03 9.626719e-04 4.223969e-04
 [8] 2.357564e-04 8.840864e-05 1.964637e-05

$density
 [1] 4.223968e-04 2.003929e-03 2.426326e-03 1.601179e-03 1.817289e-03 9.626719e-04 4.223969e-04
 [8] 2.357564e-04 8.840864e-05 1.964637e-05

$mids
 [1]  50 150 250 350 450 550 650 750 850 950

$xname
 [1] "igf1"

$equidist
 [1] TRUE
```



Box Plot

Box plot

The box plot is built using the function `boxplot()`.

This graphical representation is suitable for discrete and continuous quantitative data.

The central line of the rectangle (box) represents the median of the observations. The lower and upper extremes of the box represent the first and third quartiles, respectively.

The dashes at the end of the vertical lines can represent:

- 1) the minimum and maximum of the sample

```
> boxplot(igf1,range=0)
```

- 2) the lowest and the highest values of the sample, which are not considered outliers

```
> boxplot(igf1)
```





There are some arguments available for this representation:

- color the box make an horizontal representation

```
> boxplot(igf1,horizontal=T,col="magenta")
```

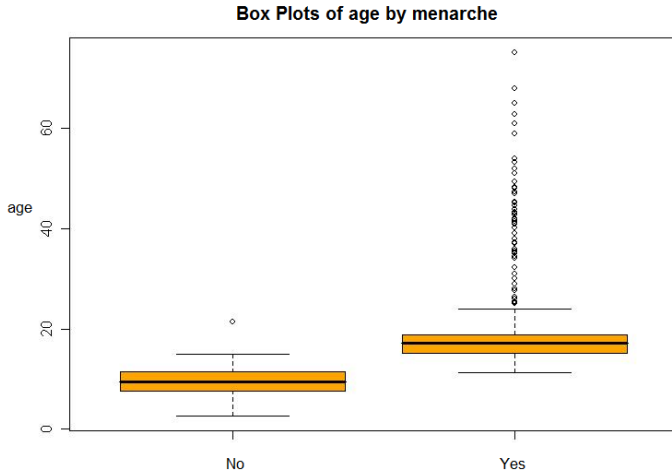
- insert title and subtitle

```
> boxplot(igf1,col="pink",main="IGF-1",sub=paste("Total:",  
+ length(igf1),"\n","NA:",sum(is.na(igf1))))
```



Box plots can also be made by groups:

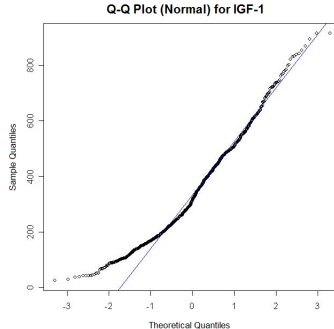
```
> boxplot(age~menarche,col="orange",main="Box Plot  
+ of age, per menarche (No/Yes)",ylab="AGE")
```



Quantile-Quantile Plot

Through the function `qqnorm()` we can represent the empirical quantiles (sample) vs. the theoretical quantiles according to the normal distribution. If the observations 'follow' a normal distribution, the points shall be provided on a line.

```
> qqnorm(igf1,main="Q-QPlot to IGF-1")  
> qqline(igf1,col="blue")
```



Acknowledgements:

Carina Silva-Fortes (ESTeSL – IPL), for allowing the use of some material produced by both of us in previous courses.

Bibliography:

Keen, K. J. (2010). *Graphics for Statistics and Data Analysis with R*. Chapman & Hall.

Maindonald, J. (2010). *Data Analysis and Graphics Using R: an Example-Based Approach*. Cambridge University Press.